

Summary

Recent M.S. Data Science graduate with hands-on experience extracting insights and building predictive models across NLP, recommender systems, and large-scale data pipelines using Python, PyTorch, scikit-learn, and Spark.

Education

M.S. in Data Science <i>University of California, San Diego</i> University GPA: 3.80 /4.00	2024 - 2026 La Jolla, CA
B.S. in Management Information Systems <i>California State University, Sacramento</i> University GPA: 3.80 /4.00 (Magna Cum Laude)	2022 - 2024 Sacramento, CA
A.S. in Accounting A.S. in Business Administration	2020 - 2022

Projects

Aspect-Based Sentiment Analysis of Google Map Reviews (2026) Compared classical NLP, transformer, and LLM approaches for aspect-level sentiment extraction across 22M+ real-world Google reviews. - Engineered 14 multi-label targets via weak supervision using star ratings as proxy labels (7 aspects × 2 sentiments) across 70.6M raw CA reviews (25GB+) - Benchmarked TF-IDF + Logistic Regression baseline with negation handling and class weighting against fine-tuned RoBERTa and GPT-4.1-mini zero/few-shot prompting - Found GPT-4.1-mini zero-shot outperformed all trained models on manually labeled ground truth (Weighted F1 0.757), RoBERTa led on weakly labeled test set (Weighted F1 0.930), revealing noise inherent in weak supervision	(Capstone Project, Paper , GitHub)
Amazon Sports & Outdoor Products Recommender System (2025) Hybrid recommender system designed to surface relevant products from a sparse 4M+ Amazon review dataset. - Mitigated 99.998% sparsity by combining collaborative filtering (SVD) with content-based features (Ridge regression) to improve cold-start performance - Benchmarked five recommendation strategies and optimized hybrid weighting (60% SVD, 40% Ridge) to balance personalization and coverage - Achieved strong ranking quality (nDCG@10 0.9981) with 31.5% RMSE reduction over baseline	(Academic Project, GitHub)
NYC Yellow Taxi Fare Prediction (2025) Fare prediction model built from 150M+ NYC taxi records, earned 3rd place in class competition. - Engineered time/efficiency features from 20GB+ trip data using Apache Spark and Dask - Trained regressor achieving MAE \$1.91 and R² 0.94, improving RMSE by 49% over baseline - Scaled training with distributed computing (15 workers, 60 cores) and validated performance using residual analysis	(Group Project, My Contributions GitHub)

Skills

Programming & Libraries: Python (Pandas, NumPy), SQL, JavaScript (D3.js)
ML & NLP: Scikit-learn, Transformers (HuggingFace), PyTorch, Regression, Classification, Recommendation Systems, Hyperparameter Tuning
Data Engineering: ETL Pipelines, Apache Spark, Dask, Large-scale Data Processing, APIs
Tools & Platforms: Git, Azure, Tableau, Excel

Experience

DBI Beverage <i>Lead AM Warehouse Associate</i>	2014 - 2020 Stockton, CA
- Managed 7-person team distributing 50K+ cases/day for Coors, Modelo, and Corona brands - Spearheaded new Warehouse Management Software (VIP EasyOps) implementation, modernized workflows - Improved inventory accuracy 12% through better labeling, FIFO procedures, and discrepancy investigation	